1 Limitations

The axiom view, and SAI as its formal model, leave out much of what scientific theories are and do. For each limitation we say whether it is partially addressable, and by what, or a genuine limit of a declarative view. The judgements are arguments, except where a theorem is cited.

Statistical confrontation with data. In physics a theory meets data statistically: with instrument and noise models it assigns likelihoods to data, often near-raw (Section 6.3), and statistical theories answer with probabilities. The proof-requiring core of SAI has no native noise model: a datum is proved or not. Graded fit enters only through randomised theories in M_A (Propositions 9.3 and G.3), through the generative variant, whose updating we prove given a support condition and whose identification we only sketch (Proposition A.23 and Remark A.24), or through sentences about probabilities, whose relation to frequencies is not deductive. It matters: on noise-free interval data the syntax of the law texts decides among the forms compatible with the Mars data, while with graded fit the data decide (Table 6 and Example 10.6). So the proof-requiring core is the noise-free special case; its natural extension inside M_A is a law with a calibrated error model; a native treatment is open.

Written theories and their readers. A written theory joined to a simple reasoner does not solve the problems the theory is used for; we see three gaps. *The bridge gap:* a problem posed in ordinary terms becomes formal only through idealisations and bridge principles that the theory does not contain (cf. Cartwright, 1983; Morgan and Morrison, 1999; an extractor prices part of the bridge, Proposition 6.4). *The computational gap:* theory and mathematics suffice logically, but finding the answer needs methods and lemmas (Section 10.5). In the model of Section 7 this gap is a theorem, visible only when what is written, or the time to write it, is charged: program length sees neither reasoner nor vocabulary (Theorem 7.2), written size depends exponentially on both (Corollary 7.10), and charged deduction cannot express the gap (Theorem E.26; derivation sizes sketched in Remark E.27). So insufficiency for a simple reasoner is a property of the pair (theory, reasoner) and of the vocabulary, in the model of Section 7, when what is written, or the time to write it, is charged; this is not a theorem about textbooks, nor about what people carry in their heads (Section 7.9). *The tacit gap:* skills (Polanyi, 1958, 1966) and exemplars (Kuhn, 1970, 1977) are not sentences, and in SAI exemplars are only data: a genuine limit.

Approximate truth. "Kepler's laws are only approximately true" is a limitation of proof-requiring compatibility with deterministic theories, not of the axiom view as SAI formalises it: its mixture contains randomised theories, and under its pricing an approximately true law is a law with a calibrated error model (Proposition 9.3). The error model is carried by the randomness of the theory machine, not stated as an axiom; a treatment in which it is stated, by sentences about probabilities, belongs to the open problem stated above. Whether a deterministic theory holds a law honestly or with exceptions is a modelling choice of pricing (Proposition G.16); humility is partiality for deterministic theories, calibration for randomised ones (Proposition G.22). A canonical metric for approximate predictions, which we think does not exist, and a general comparative notion of truthlikeness (cf. Popper, 1963; Miller, 1974; Tichý, 1974; Niiniluoto, 1987) are open; our notion is local: a question family, a tolerance, a domain.

Domains, several models, and inconsistent theories in use. Practice uses mutually inconsistent models of one system, each in its domain (the Earth as flat, as a sphere, as an oblate

feature of concepts	axiom view	how captured, or what is missing
for questions and problems, not predictions	yes/no	yes for questions posed in a language (§ 6.3); no for problems needing search or action
inert alone; parts of a mind	a limit	worth depends on reasoner and vocabulary (Cor. 7.10); free deduction hides it
inferential role constitutive	yes	axioms are inferential relations (meaning postulates, Ramsey–Lewis)
compositional	yes/no	yes for contents, as formulas; no for recognizers, which are procedures
learning by new links	partly	a persisting symbol, but a global mixture over whole theories
soft; meaning change	partly	"within ε of a sphere" is a true partial theory; not meaning change
analogy	partly	partial interpretations, rewarded as they compress (1); not finding them
new concepts from old	partly	definitions cost $O(1)$ (Cor. 3.10), help a chaining reasoner (Prop. 7.8)
methods attached	partly	lemmas for a bounded reasoner, or axioms via Craig at exponential cost (Lem. F.2); not motor methods
recognizers	no	a procedure on situations; the extractor of Prop. 6.4, taken as given
one notion, many recognizers	yes	identity is the symbol with its axioms; recognizers are bridges
built from pointers	yes	a constant fixed by one axiom, which gathers further axioms
purposes	no	axioms can state functions; their use in search is not modelled
clarification; facing action	no	not in the inducer
natural types	partly	sorts and arities; their emergence is not modelled

Table 1: Features of human concepts, and whether the axiom view captures them (an assessment).

spheroid), and proof-requiring SAI gives an inconsistent theory weight 0. A sound device exists: honest forms "q lies within ε of f on the domain" of true models are jointly consistent over a true background; exact ones can prove everything (Section 9.4). But physicists rarely state or propagate worst-case bounds on the error of an approximate model (measurement uncertainties are propagated routinely, and statistically); worst-case propagation becomes vacuous over long chains; and Lanford's theorem, the rigorous honest form of the Boltzmann equation as described in expositions (Lanford, 1975), holds only for a fraction of the mean free time in the low-density limit, while practice rightly uses the equation far beyond. How to choose a model, a domain or a tolerance is left to tacit understanding; under SAI's pricing a domain is usually a tolerance with an error distribution varying with the system (Section 9.7), and knowing when a model applies is a recognizer, a genuine limit.

Concepts. Table 1 lists features of human concepts that an account of concepts as pieces of code, combined to generate predictions, misses, and whether the axiom view captures them.[1] We accept them as features of human concepts, the first only comparatively (Section 6.1), with two further qualifications: inferential role is one major account of meaning (cf. Sellars, 1953), disputed by referential accounts (cf. Kripke, 1980); and natural types we find plausible. One feature supports the axiom view: a notion with different recognizers has its identity fixed otherwise.

For analogy there is a bound: with $K(T)$ the least $K(e)$ such that $\mathsf{W}_e = T$, for every r.e. theory T_1 and computable map τ from sentences to sentences

$$K\big(T_1 \cup \tau[T_1]\big) \;\leq\; K(T_1) + K(\tau) + c, \tag{1}$$

since an index of $T_1 \cup \tau[T_1]$ is computable from indices of T_1 and τ, and $K(x,y) \leq K(x) + K(y) + c$. Structure mapping (Gentner, 1983) and models in science (cf. Hesse, 1963) account for analogy;

[1] The list follows, with changes, an unpublished draft essay of Hänni's on concepts.

patchy or stretching concepts (cf. Wilson, 2006; Lakatos, 1976) and interpretation as abduction (Hobbs et al., 1993) are outside the view.

Other limitations. *Science predicts very particular things, in set-ups it builds.* SAI captures half, with selected data and partial theories; which questions are worth asking (value-weighted questions would be a handle), and the making of phenomena (cf. Hacking, 1983), are outside it. *Concept formation and discovery.* Step 3 of the ladder (Section 6.1) is partially addressable (Theorem 8.4 and Corollary 7.10); step 4, the revision of the base abstractions, is a genuine limit for now; and SAI evaluates theories in a fixed space of hypotheses without discovering them. *Theory-laden and fallible data.* A base language may presuppose a theory (heliocentric longitudes presuppose Copernican geometry; Definition 10.1; Section 6.1), which biases comparisons between theories that disagree about it; and a deterministic theory must prove even a mislabelled datum, by an exception axiom; a noise component in the likelihood, as in our earlier report (Hänni and Claude (Anthropic), 2026), is a partial remedy. *Purposes and technology* are outside the view, and *science as a social structure* is outside a model with one inducer: genuine limits, the last perhaps where Question 6.1 finds its answer. *The mathematics is free, and large.* Building a background B' into B gains at most $K(B') + c$ bits on any data (Proposition 3.8), a constant that does not grow with the data but is not small for the mathematics physics uses, and that matters when rival theories need different mathematics. More generally, all our comparisons of theories are relative to the universal machine, up to constants that are not small at the scale of real theories (Proposition 3.7). With free deduction, data the background entails carry no information (Proposition 5.15), so induction in mathematics needs a reasoner for whom deduction is costly.

What the axiom view captures. It captures the static, declarative and inferential side of theories and concepts (contents, inferential relations and their composition, learning by weighing axioms globally, pointers, partiality, approximate truth in honest or error-model form) and misses the procedural, the statistical and the dynamic side. We offer a model of the declarative core of scientific theories, in which the procedural side reappears formally as the dependence of a written theory's worth on its reasoner and its vocabulary (Section 7).

References

Nancy Cartwright. *How the Laws of Physics Lie*. Oxford University Press, Oxford, 1983.

Dedre Gentner. Structure-mapping: A theoretical framework for analogy. *Cognitive Science*, 7 (2):155–170, 1983.

Ian Hacking. *Representing and Intervening: Introductory Topics in the Philosophy of Natural Science*. Cambridge University Press, Cambridge, 1983.

Kaarel Hänni and Claude (Anthropic). Bayesian axiom induction from instances: Template priors, derivation likelihoods and time penalties. Report, 2026. Folder `axiom-induction` of `https://github.com/kaarelh/AI-works`.

Mary B. Hesse. *Models and Analogies in Science*. Sheed and Ward, London, 1963.

Jerry R. Hobbs, Mark E. Stickel, Douglas E. Appelt, and Paul Martin. Interpretation as abduction. *Artificial Intelligence*, 63(1–2):69–142, 1993.

Saul A. Kripke. *Naming and Necessity*. Harvard University Press, Cambridge, MA, 1980.

Thomas S. Kuhn. *The Structure of Scientific Revolutions*. University of Chicago Press, Chicago, 2 edition, 1970.

Thomas S. Kuhn. Second thoughts on paradigms. In *The Essential Tension: Selected Studies in Scientific Tradition and Change*, pages 293–319. University of Chicago Press, Chicago, 1977.

Imre Lakatos. *Proofs and Refutations: The Logic of Mathematical Discovery*. Cambridge University Press, Cambridge, 1976. Edited by John Worrall and Elie Zahar.

O. E. Lanford, III. Time evolution of large classical systems. In J. Moser, editor, *Dynamical Systems, Theory and Applications*, volume 38 of *Lecture Notes in Physics*, pages 1–111. Springer, Berlin, 1975.

David Miller. Popper's qualitative theory of verisimilitude. *British Journal for the Philosophy of Science*, 25(2):166–177, 1974. doi: 10.1093/bjps/25.2.166.

Mary S. Morgan and Margaret Morrison, editors. *Models as Mediators: Perspectives on Natural and Social Science*, volume 52 of *Ideas in Context*. Cambridge University Press, Cambridge, 1999.

Ilkka Niiniluoto. *Truthlikeness*, volume 185 of *Synthese Library*. D. Reidel, Dordrecht, 1987.

Michael Polanyi. *Personal Knowledge: Towards a Post-Critical Philosophy*. University of Chicago Press, Chicago, 1958.

Michael Polanyi. *The Tacit Dimension*. Doubleday, Garden City, NY, 1966.

Karl R. Popper. *Conjectures and Refutations: The Growth of Scientific Knowledge*. Routledge and Kegan Paul, London, 1963.

Wilfrid Sellars. Inference and meaning. *Mind*, 62(247):313–338, 1953.

Pavel Tichý. On Popper's definitions of verisimilitude. *British Journal for the Philosophy of Science*, 25(2):155–160, 1974. doi: 10.1093/bjps/25.2.155.

Mark Wilson. *Wandering Significance: An Essay on Conceptual Behavior*. Oxford University Press, Oxford, 2006.